TicketMastery: concert ticketing on AWS

IS458 Cloud Management and Engineering, final report

Cody, Andrew, Nai Xin, Yan Ting, Shi Hui, Elroy

10 November 2026

1. Executive summary

TicketMastery is a Singapore-based B2C platform where fans across Asia-Pacific buy concert tickets. It falls under the course's *Online Shopping System* theme. Ticket sales are spiky: a popular on-sale can bring 50,000–100,000 concurrent fans in its first minutes, about 10× normal load. The platform must stay up during that spike and must never sell a seat twice.

The stack runs in ap-southeast-1 (Singapore):

- **Front end:** CloudFront and WAF serve the React single-page app from S3 and forward `/api/*` to an API Gateway HTTP API.
- **API:** API Gateway checks Cognito tokens and reaches an internal Application Load Balancer through a VPC link. ALB path rules send each use case to its own ECS Fargate service (six services from one image, including a read-only reporting service), each scaling on its own. Concerts, seats and orders are stored in RDS PostgreSQL.

Two event-driven parts sit beside the request path:

- **Virtual waiting room:** SQS FIFO, DynamoDB, and the `queue-manager` Lambda run every minute by EventBridge Scheduler. It admits 500 fans per concert per minute, so the spike never reaches the database at once.
- **E-ticket pipeline:** SNS → SQS → the `ticket-fn` Lambda → S3 and SES. It keeps PDF rendering and email off the checkout path.

No seat is ever sold twice, because a seat hold is one all-or-nothing SQL `UPDATE` in the same database as the seats and orders.

The report presents two designs, both built from the same Terraform; a block of variables switches between them:

- **Target high-availability (HA) design:** two AZs, private subnets with a NAT gateway per AZ, at least 2 tasks per service, Multi-AZ RDS; ≈ USD 310/month for the whole stack.
- **Implementation mode:** what the instructor asked us to deploy at the proposal review. One AZ, no NAT gateway, one task per service, single-AZ RDS; ≈ USD 150/month for the whole stack (≈ 136 of it the base stack).

Section 3 shows both diagrams. Section 7 says which availability targets are therefore design-only.

Operations are managed as code:

- Terraform modules with shared S3 remote state and locking;
- GitHub Actions CI on every pull request and push to `main` (API type-check and tests against Postgres 16 with a coverage gate, `ticket-fn` tests, `terraform fmt/validate`), and a frontend deploy over OIDC;
- review routed by CODEOWNERS;
- 31 CloudWatch alarms and a dashboard, alerting by email through SNS;
- AWS Budgets at USD 20/40/60/80;
- k6 tests and failure drills for every NFR.

2. Business context and requirements

2.1 Context, users and goal

TicketMastery has three kinds of user (Google Doc, *Customer Base*):

- **Guests** browse concerts without signing in.
- **Fans** (registered users) join waiting rooms, buy tickets and download e-tickets.
- **Organisers** (internal staff, Cognito group `organisers`) create concerts, ticket types and seat layouts, and watch sales live.

The business goal (Google Doc): fans buy tickets *"without the platform buckling under traffic spikes, with every seat sold to exactly one buyer"*. Checkout stays secure, e-tickets arrive within seconds, and organisers see pricing, inventory and sales in real time.

2.2 Base requirement and traffic model

The specification's base requirement for this theme: *the company covers Asia-Pacific; 80% of its customers are from Singapore; the remainder are split evenly between Malaysia, Thailand and Hong Kong.*

During the proposal we drafted a different split (Google Doc, *Traffic model*): Singapore 60%, Malaysia 20%, Indonesia/ Thailand/Philippines 15%, rest of APAC 5%. It was labelled "team assumption, to confirm". **This report uses the specification's split**, because that is the base requirement we are assessed against.

Origin	Share	How it is served
Singapore	80%	CloudFront edge in Singapore; API in ap-southeast-1
Malaysia, Thailand, Hong Kong	≈ 6.7% each	CloudFront edge in-country; dynamic calls cross the AWS backbone to Singapore

The design conclusion is the same under either split: one region plus CloudFront edges is enough. Our earlier serverless write-up estimated round trips from Singapore of about 10 ms to Kuala Lumpur and 30–40 ms to Bangkok and Hong Kong, well inside the 300 ms p95 target. The distribution uses `PriceClass_200`, which has edge locations in all four markets.

Load assumptions for one major on-sale (Google Doc):

Phase	Concurrent users	Absorbed by
Normal day	≈ 5,000	CloudFront for static content; ALB and the minimum number of tasks for the API
First 10 minutes of a major on-sale	50,000–100,000	The waiting room queues them, CloudFront serves the waiting page, and 500 per concert are admitted each minute

Capacity math. 100,000 waiting fans polling `GET /status` every 5 s is ≈ 20,000 req/s, and `/join` alone at the 100 rps stage throttle takes ≈ 1,000 s for 100,000 fans. Today the excess gets `429`; a fan keeps their place in DynamoDB and the page simply polls again. For production: per-route throttles, `429` + `Retry-After` with jittered backoff in the page, an API Gateway quota raised before each on-sale, and a public per-concert "admitted up to position N" value cached at CloudFront (≈ 5 s), so the page works out its own place without a per-fan request. NFR-1's 10 req/s baseline assumes the ≈ 5,000 normal-day users are mostly served static content by CloudFront; the 10× test (100 req/s) is capped by that same throttle, so it proves the stack up to the throttle.

2.3 Use cases and functional requirements

The proposal listed ten fan, two admin and six system use cases. We grouped them into five deliverable use cases. Each is one API module and ECS service (UC5 adds a read-only reporting service) and one frontend folder, with one owner:

UC	Name	Owner	Functional requirements
UC1	Concert discovery	Cody, Nai Xin	FR-1 list published concerts, with search (`q`, artist, venue) and date filters; detail with ticket types, prices and availability. FR-2 seat map showing available, held and sold seats (an expired hold shows as available)
UC2	My Tickets & Account	Andrew	FR-3 sign up, sign in and sign out with Cognito (email verification). FR-4 list the fan's confirmed orders, open an e-ticket, download its PDF through a short-lived link
UC3	Virtual waiting room	Yan Ting	FR-5 join, see the queue position, be admitted in arrival order in batches. FR-6 only admitted fans may hold seats; admission lasts 15 minutes
UC4	Seat selection & checkout	Shi Hui	FR-7 hold 1–3 seats for 10 minutes, all or nothing, one active hold per fan per concert. FR-8 simulated payment: success sells the seats, failure or timeout releases them. FR-9 email the e-ticket PDF after payment
UC5	Concert & ticket management	Elroy	FR-10 create, edit, publish and cancel concerts; ticket types and seat layout, locked once the sale opens; cancelling emails every ticket holder. FR-11 live available, held and sold counts and revenue. FR-12 sales reporting (proposal: for a BI tool)

The six system use cases (autoscaling, flash-sale queue, monitoring, caching, load balancing, failure handling) are covered in Sections 7 and 8. Payment is simulated: the proposal scoped real payment as design-only.

2.4 Non-functional requirements

These are the proposal's NFRs (Google Doc). Section 10 has the method and the result for each one.

#	Category	Target
NFR-1	Performance	Browse and seat-map APIs: p95 < 300 ms, p99 < 800 ms at baseline
NFR-2	Scalability	Sustain 10× baseline; scale-out within 5 min; 0 dropped requests at the admitted rate
NFR-3	Correctness	0 oversold seats across all tests
NFR-4	Hold accuracy	Seat available again within 30 s of hold expiry
NFR-5	Availability	99.9% monthly for the purchase path; no outage on the loss of one instance or one AZ
NFR-6	Fairness	Admission order matches queue order
NFR-7	Delivery	95% of e-ticket emails sent within 60 s of payment
NFR-8	RPO	Paid orders: 0 on AZ loss; ≤ 30 min on data corruption; ≤ 24 h on region loss
NFR-9	RTO	AZ loss ≤ 5 min; data corruption ≤ 1 h; region loss ≤ 4 h (design only)
NFR-10	Security	HTTPS on public paths; encryption at rest; DB only in private subnets
NFR-11	Cost	Total project spend under USD 80

3. Architecture

3.1 From the serverless proposal to containers and RDS

Our first write-up was all-serverless: API Gateway REST and WebSocket APIs, one Lambda per module, and DynamoDB for concerts, seats and orders. Much of its reasoning still holds, and we kept it: one region with CloudFront for APAC; WAF on the single public entry point; SQS FIFO for the queue; SNS/SQS fan-out to idempotent consumers; Cognito for sign-in.

We moved the core to **ECS Fargate + RDS PostgreSQL** for four reasons.

1. **Correctness is easier to prove in SQL.** A hold is one conditional `UPDATE` with a row-count check. Confirming, cancelling and locking the layout are each one transaction, guarded by foreign keys, check constraints and triggers. In DynamoDB the same guarantees need careful single-table design, `TransactWriteItems` is capped at 100 items, and reporting needs exports.
2. **Organiser reporting is relational.** Sales by concert, ticket type and day are SQL views (Section 6.5).
3. **It covers the course's infrastructure topics.** The serverless design had nothing to put in a VPC, so it showed no subnets, security-group chains, load balancer, container autoscaling or Multi-AZ failover.
4. **The cost is predictable at our scale**, and the API is TypeScript like the frontend.

We kept serverless where it fits best:

- **The waiting room**, because DynamoDB absorbs the spike without loading the single database; it was also the instructor's suggestion.
- **The asynchronous e-ticket worker.**

The trade-off is that containers and RDS cost money while idle and scale more slowly than Lambda. The waiting room makes that acceptable, because it caps how many buyers reach the API.

3.2 Target HA design

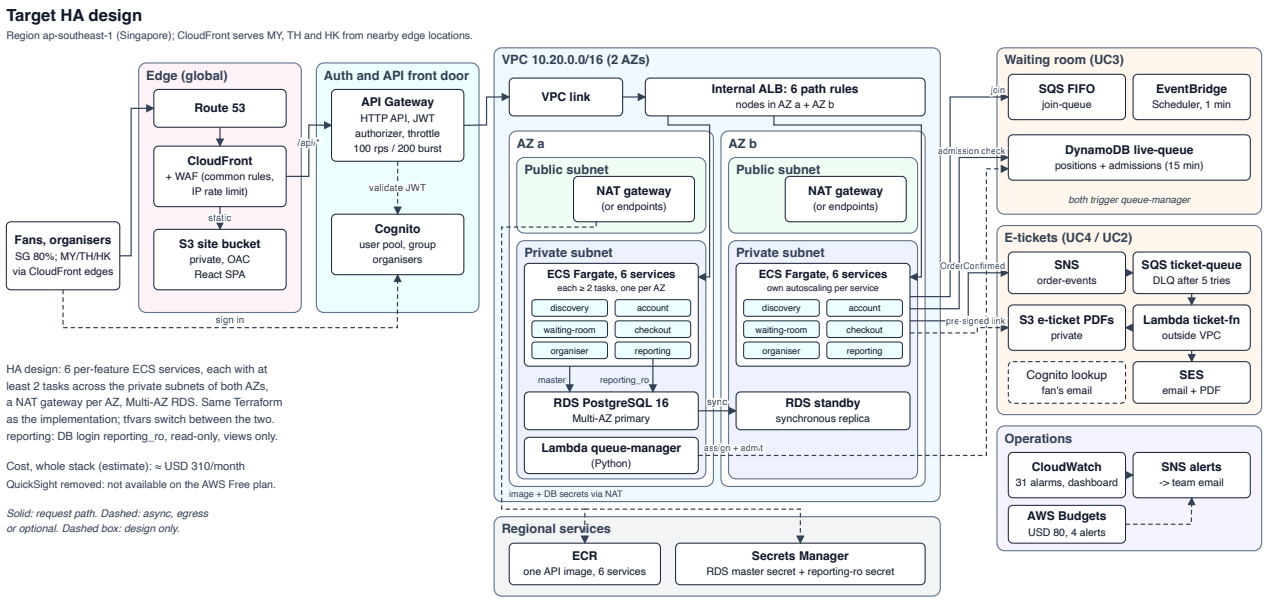


Figure 1: Target HA architecture (design; drawn by `figures/architecture.py`).

- **Edge.** Route 53 → CloudFront (`PriceClass_200`, redirect to HTTPS) with a WAF web ACL: the AWS managed common rule set plus an IP rate rule at 2,000 requests per 5 minutes. The SPA comes from a private S3 bucket through Origin Access Control (OAC). `/api/*` and `/health` go to API Gateway.
- **API front door.** API Gateway HTTP API.
  - `GET /api/concerts…` and `GET /health` are public. Every other `/api/*` route, the waiting room included, needs a Cognito access token (JWT authorizer).
  - The stage throttles at 100 rps steady and 200 burst, and writes JSON access logs.
  - A **VPC link** reaches an **internal** ALB, so the ALB is never internet-facing.
- **Compute.** Six ECS Fargate services from one image (Node 20, Express, TypeScript), one per use case plus reporting: each at least two tasks, one per AZ, in private subnets, with its own CPU and request-count autoscaling. Egress goes through a NAT gateway per AZ (or interface endpoints).
- **Data.** RDS PostgreSQL 16 Multi-AZ (synchronous standby, automatic failover), 7-day point-in-time recovery, deletion protection, encrypted storage. RDS manages the master password in Secrets Manager.
- **Waiting room.** SQS FIFO `join-queue` → `queue-manager` (Python 3.13, in the VPC) → DynamoDB `live-queue`. EventBridge Scheduler (`rate(1 minute)`) runs the admit step.
- **E-tickets.** SNS `order-events` → SQS `ticket-queue` (DLQ after 5 receives) → `ticket-fn` (Node.js 22, arm64, outside the VPC) → S3 `tickets/<orderId>.pdf` + SES.
- **Operations.** CloudWatch alarms and a dashboard → SNS `alerts` → email. AWS Budgets. An RDS read replica for reporting at production scale (design only).

3.3 Implementation mode (deployed)

At the proposal review the instructor asked us to **skip the NAT gateway, use one AZ and not implement HA, but still show the HA design**. The Terraform defaults build exactly that (Figure 2).

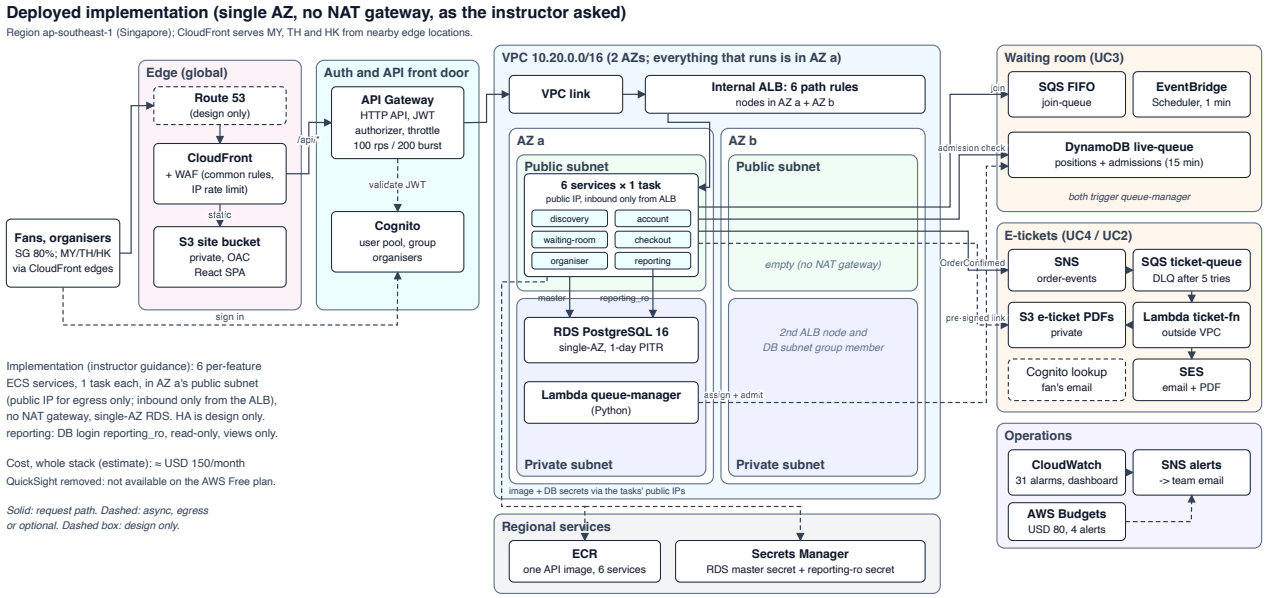


Figure 2: Deployed implementation (drawn by `figures/architecture.py`).

Aspect	Implementation (deployed)	Target HA	Variable(s)
Task subnets	AZ a public subnet; public IP used only for egress	Private subnets, AZ a + b	`tasks_in_public_subnets`, `tasks_single_az`
Task egress	Own public IP via the internet gateway	NAT per AZ or interface endpoints	`enable_nat_gateway`, `single_nat_gateway`, `enable_vpc_endpoints`
Tasks / scaling	6 services × 1 task, autoscaling per service (max 2–3)	6 services × ≥ 2 tasks, autoscaling per service	`ecs_services`, `desired_count`, `enable_autoscaling`, `autoscaling_*_capacity`
Database	Single-AZ, pinned to AZ a, 1-day backups	Multi-AZ, 7-day backups, deletion protection	`db_multi_az`, `db_backup_retention_period`, `db_deletion_protection`
Cost (estimate, whole stack)	≈ USD 150/month	≈ USD 310/month	`docs/cost-estimate.md`

- **The tasks' public IPs are not an open door.** The task security group has one ingress rule: the container port, from the ALB security group. RDS stays private (`publicly_accessible = false`).
- **Why there are still two subnets of each kind.** An ALB needs subnets in two AZs, and an RDS subnet group must span two AZs. In implementation mode AZ b holds only the second ALB node and an empty member of the DB subnet group.
- **Lambdas and the VPC.** A Lambda attached to a VPC never gets a public IP, so with no NAT it reaches only what the free S3 and DynamoDB gateway endpoints cover.
  - `ticket-fn` needs SQS, S3, SES and Cognito, so it runs **outside** the VPC.
  - `queue-manager` runs in the private subnets. Its SQS input arrives through an event source mapping (the Lambda service polls the queue), it writes DynamoDB through the gateway endpoint, and it publishes its metric as a log line in CloudWatch Embedded Metric Format (EMF), because it cannot call the CloudWatch API.

4. Resource allocation and justification

The table sizes the implementation mode. Prices are on-demand list prices for ap-southeast-1 at 730 h/month, from `docs/cost-estimate.md` and the module READMEs. They are estimates, not bills.

Layer	Service and configuration	Justification	≈ USD/month
Edge	CloudFront `PriceClass_200`; SPA cached 1 day (hashed assets 1 year)	Cheapest price class with edges in SG, MY, TH and HK; TLS close to fans; absorbs browse traffic	per request
Edge security	WAF: AWS common rules + IP rate limit 2,000 / 5 min (2 rules)	Blocks exploits and single-source floods before they reach the API	≈ 7 (5 per ACL + 1 per rule)
Static site	Private S3 + OAC	No servers for static content	< 1
Identity	Cognito user pool, `organisers` group, email verification	Managed sign-in; JWTs that API Gateway validates	free tier
API front door	API Gateway HTTP API, JWT authorizer, VPC link, throttle 100 rps / 200 burst	Central API management (instructor feedback a). Auth before compute. An HTTP API's VPC link targets an ALB directly; a REST API would need an extra NLB (≈ 18+). The throttle is sized to protect one 0.25 vCPU task	per request
Load balancer	Internal ALB, health check `/health` (runs DB `SELECT 1`)	Spreads requests over tasks and AZs; takes unhealthy tasks out	24.2
Compute	Fargate, 6 services × 1 task, each 0.25 vCPU / 0.5 GB, x86_64; per service a target group, ALB path rule and autoscaling (CPU 60%, 500 requests per task per minute; max 2–3)	No servers to patch; per-second billing. Each use case scales and deploys on its own (feedback b). Sizing: one task served ≈ 60 req/s (NFR-2), so 3 tasks cover the 100 rps throttle for a hot service. x86 because the image is built for amd64 (ARM64 ≈ 20% cheaper: future). No Fargate Spot: an interrupted task mid-on-sale costs sales (Spot is possible for reporting)	6 × 11.2 = 67.2
Public IPv4	1 per task, 6 (the ALB is internal, so it has none)	Replaces a NAT gateway (≈ 43 per AZ)	6 × 3.65 = 21.9
Database	RDS PostgreSQL 16 `db.t4g.micro`, 20 GB gp3, encrypted, single-AZ	Transactions for zero double-booking; relational reporting. Smallest Graviton class	18.3 + 2.8
Secrets, registry, logs	Secrets Manager (RDS-managed master secret; `reporting_ro` secret); ECR (scan on push, keep 10 images); CloudWatch Logs (14 days)	No password in code or state; private images; central logs	0.8 + 0.1 + 0.9 = 1.8
Waiting room	SQS FIFO `join-queue` (DLQ after 3 tries); DynamoDB `live-queue` (on-demand, TTL); `queue-manager` (128 MB, 30 s, SQS batch 10); Scheduler `rate(1 minute)`	FIFO keeps arrival order per concert. DynamoDB scales with the spike without loading RDS. 43,800 runs/month fit the Scheduler (14 M) and Lambda (1 M) free tiers	≈ 0
E-tickets	SNS → SQS `ticket-queue` (visibility 6 × the 60 s timeout, DLQ after 5, DLQ kept 14 days) → `ticket-fn` (Node.js 22 arm64, 256 MB, batch 10, max concurrency 2) → private S3 + SES	Checkout doesn't wait for the PDF or the email. Concurrency 2 respects the SES sandbox rate of 1 email/s	≈ 0
Monitoring	31 alarms (44 alarm metrics); 1 dashboard (19 graphs, 63 metrics); 4 custom metrics; SNS email; Budgets	Section 8	≈ 6.4 (34 billable alarm metrics × 0.10 + 3 for a dashboard over 50 metrics)
Whole stack			**≈ 150** (136.2 base stack + WAF 7 + monitoring 6.4 + pay-per-request ≈ 0.3 = 149.9; ≈ 4.9/day)

The target HA design is **≈ USD 310/month** (≈ 10.2/day; Appendix B). The extra cost is two NAT gateways with their Elastic IPs (≈ 94), Multi-AZ RDS (≈ +21) and a second task per service (6 × 11.2 ≈ +67), minus the six public IPs (≈ −22).

Alternatives we rejected:

- **A NAT gateway in implementation mode.** ≈ 47/month (NAT + Elastic IP) for egress that six task IPs give for ≈ 22.
- **ElastiCache for seat holds** (proposal). A hold in Redis and an order in RDS are two systems that can disagree; one SQL `UPDATE` cannot.

- **QuickSight.** The account is on the AWS Free plan, which can't subscribe to it; a reporting service replaces it (Section 6.5).

5. Linkage of resources

5.1 Request path per use case

UC	Path
Static	Browser → CloudFront (+ WAF) → S3 via OAC; a CloudFront function rewrites SPA routes to `index.html`
UC1	CloudFront `/api/concerts…` → API Gateway public route → VPC link → internal ALB (one path rule and target group per service; unmatched paths and `/health` go to `discovery`) → Fargate `discovery` → RDS
UC2	Cognito sign-in (access token) → CloudFront → API Gateway (JWT) → ALB → `account` → RDS. Download: the API sends `HeadObject` to S3 for `tickets/<orderId>.pdf` and returns a 300 s pre-signed `GetObject` URL; the browser downloads straight from S3
UC3	`POST /join` → `waiting-room` → SQS FIFO (`MessageGroupId` = concert) → `queue-manager` → DynamoDB. `GET /status` → `waiting-room` → DynamoDB `GetItem`. Scheduler → `queue-manager` admit step → DynamoDB
UC4	`POST /api/checkout/holds` → `checkout` → DynamoDB `GetItem` (admitted?) → RDS transaction. `…/confirm` → RDS transaction → SNS `order-events` → SQS `ticket-queue` → `ticket-fn` → Cognito `AdminGetUser` (email) → S3 PDF → SES → S3 `.sent` marker
UC5	Organiser → API Gateway (JWT) → `organiser` (checks the `organisers` group) → RDS; the dashboard is polled every 5 s. Reports: `/api/reports/*` → `reporting` (DB login `reporting_ro`) → the `report_*` views

Figure 3 traces one purchase through the waiting-room and checkout services to the e-ticket email.

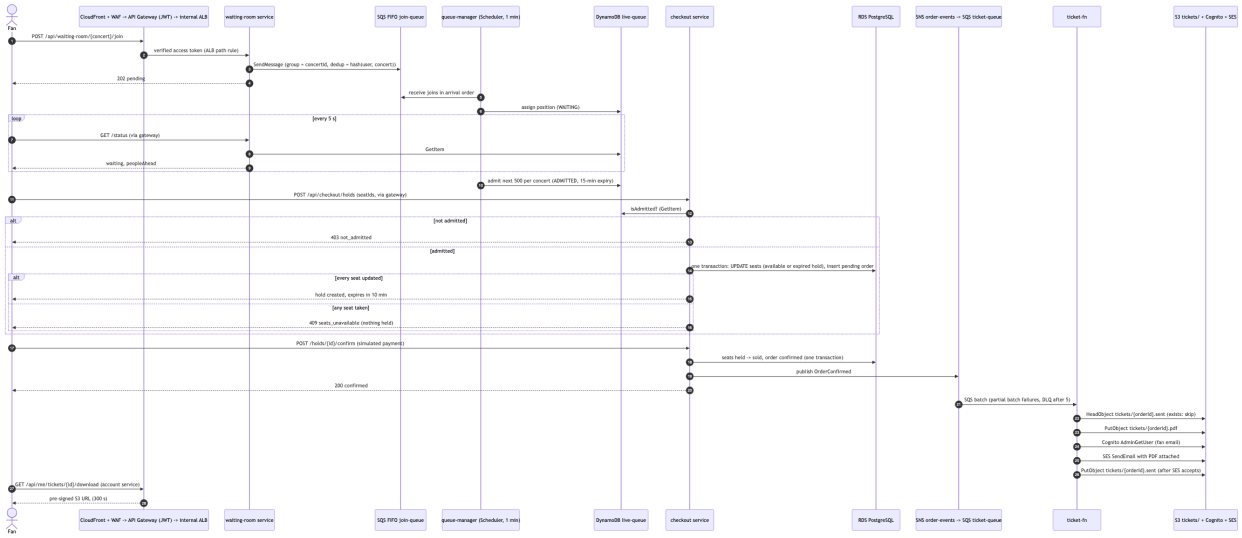


Figure 3: On-sale purchase sequence (source `figures/purchase-sequence.mmd`).

5.2 Events

Producer → channel → consumer	Message	Guarantees and failure handling
API → SQS FIFO `join-queue` → `queue-manager`	`{ eventId, userId, joinedAt, saleStartsAt, saleEndsAt }`. Group = concert; dedup ID = hash of `userId#concertId`	In order per concert; duplicates dropped. If a record fails, it and every record after it are retried, so nobody overtakes. DLQ after 3 tries (alarmed)
EventBridge Scheduler → `queue-manager`	Scheduled admit run every minute	Async, retries off (the next run is a minute away). Admits are conditional on `status = WAITING`, so overlapping runs never admit anyone twice
API → SNS `order-events` → SQS `ticket-queue` → `ticket-fn`	`OrderConfirmed` v1 from checkout; `ConcertCancelled` v1 from the organiser module, one per confirmed order (attribute `eventType`; other types are skipped)	At-least-once delivery. Partial batch failures; DLQ after 5 (alarmed); `.sent` / `.cancelled.sent` markers so redeliveries send nothing twice
`queue-manager`, `ticket-fn` → CloudWatch Logs (EMF)	`TicketMastery/WaitingRoomSize`, `TicketMastery/ETicketDeliveryLatency`	Metrics without API calls; used by alarms and the dashboard
Alarms, Budgets → SNS `alerts` / email	ALARM and OK transitions	Section 8

5.3 IAM

`queue-manager`, `ticket-fn`, the Scheduler and GitHub Actions each have their own role with least-privilege inline policies; the six API services still share one task role (a gap, Section 9.5). The module that owns a resource grants that task role access to it, so permissions sit next to the resource.

Principal	Allowed	Granted in
ECS execution role	ECR pull, its log group, read the two DB secrets (RDS master; `reporting_ro`)	`ecs-service`
API task role (shared by the six services)	`sqs:SendMessage` on `join-queue`; `dynamodb:GetItem` on `live-queue` (status and the admission check); `sns:Publish` on `order-events`; `s3:GetObject` on `tickets/*` + `s3:ListBucket`, so a missing PDF answers 404 rather than 403; `ssmmessages:*` for ECS Exec	`waiting-room`, `messaging`, `ecs-service`
`queue-manager`	Receive and delete on `join-queue`; `GetItem/PutItem/UpdateItem/Query` on `live-queue` and its indexes; VPC ENIs; logs	`waiting-room`
Scheduler role	`lambda:InvokeFunction` on `queue-manager` only	`waiting-room`
`ticket-fn`	Receive and delete on `ticket-queue`; S3 get/put on `tickets/*` + list; `ses:SendEmail` and `ses:SendRawEmail` with the condition `ses:FromAddress` = the sender; `cognito-idp:AdminGetUser` on the pool; its log group	`messaging`
GitHub Actions	OIDC, `main` branch only: list/put/delete in the frontend bucket, `CreateInvalidation` on the distribution	`github-deploy`

5.4 Network

- **Security-group chain:** VPC link SG → `alb_sg` (port 80) → `ecs_sg` (container port only) → `rds_sg` (5432 only, plus `rds_extra_ingress_sg_ids` for in-VPC consumers).
- **Routing:**
  - Public subnets route to the internet gateway.
  - Private subnets have no internet route in implementation mode; they carry the free S3 and DynamoDB gateway endpoints.
  - In the HA design, private subnets route through a NAT gateway per AZ.
- **Encryption in transit:**
  - **Encrypted:** browser → CloudFront (HTTPS, redirect from HTTP); CloudFront → API Gateway (TLS 1.2); API → RDS (`DB_SSL=true`); every AWS SDK call.
  - **Plain HTTP:** VPC link → internal ALB → task, on port 80, inside the VPC only (Section 9.5).

6. Implementation per use case and key mechanisms

6.1 Code structure

One image holds six modules (`src/modules/<uc>/`: `router.ts`, `service.ts`, `repository.ts`, `schemas.ts`); the `API_MODULES` variable picks the one each ECS service mounts.

- **Isolation.** Each use case has its own migrations (`NNNN_<uc>_<desc>.sql`), and modules don't import each other.
- **Migrations** run at task start under a Postgres advisory lock, in the five services with the master login (not in reporting). A failed migration stops the task, and the ECS deployment circuit breaker rolls back.
- **Health checks.** `/health` runs a DB query and is used by the ALB. `/health/live` skips the DB, so a database blip doesn't kill every task.
- **Lambdas** live in `lambdas/<function>/`. `ticket-fn` has node:test unit tests and a committed esbuild bundle, so `terraform plan` never runs npm; `queue-manager` has standard-library unittests.

6.2 UC3: virtual waiting room

1. **Join.** A signed-in fan calls `POST /api/waiting-room/:concertId/join`. The API sends the join to SQS FIFO (group = concert, dedup = `userId#concertId`) together with the concert's sale window, and answers `202 pending`. A fan already queued keeps their number.
2. **Number.** `queue-manager` (SQS trigger, batch 10) runs an atomic `ADD lastPosition 1` on the concert's `#COUNTER` row in `live-queue`, which also stores `saleStartsAt` and `saleEndsAt`. It then does a conditional put of the fan's row as `WAITING` with that position.
3. **Admit.** EventBridge Scheduler invokes `queue-manager` every minute. For each concert with a queue (sparse GSI `counter-index`), the run:
   - **skips** the concert unless `saleStartsAt ≤ now < saleEndsAt`, so one schedule serves concerts with different sales, and early joiners aren't admitted and timed out before tickets go on sale;
   - **admits** the next **500** `WAITING` fans by position: `ADMITTED`, `admissionExpiresAt = now + 15 min`;
   - **moves `admittedUpTo` forward**.

   Each update is conditional, so overlapping runs never double-admit. A 100-position look-back catches fans whose GSI entry arrived late.
4. **Status.** The browser polls `GET /status` every 5 s. It gets `pending`, `waiting` (with `peopleAhead = position − 1 − admittedUpTo`), `admitted` (with `admittedUntil`) or `expired`. A fan whose admission ran out can rejoin at the back.
5. **Enforcement.** Checkout's `isAdmitted` (`checkout/admission.ts`) reads the fan's row with the same mapping as the status endpoint. Only `ADMITTED` with an unexpired window may create a hold; anything else gets `403 not_admitted`. A DynamoDB error fails the request and is never treated as admitted. A hold made in the window keeps its own 10 minutes to pay.

Every admit run also emits `WaitingRoomSize` (total `lastPosition − admittedUpTo` across concerts). The batch size caps the new buyers that reach the API and RDS at 500 per concert per minute, however large the crowd. That cap is what lets one small task per service and a micro database survive an on-sale.

6.3 UC4: all-or-nothing seat hold

The hold is one statement (`src/db/seat-holds.ts`), run in the same transaction that creates the `pending` order:

```
UPDATE seats SET status = 'held', held_by = $3, hold_id = $4,
                 held_at = now(), hold_expires_at = now() + make_interval(mins => 10)
 WHERE concert_id = $1 AND id = ANY($2::uuid[])
   AND (status = 'available' OR (status = 'held' AND hold_expires_at < now()))
RETURNING id;
```

- **All or nothing.** If fewer rows come back than seats were requested, the transaction rolls back and the API answers `409 seats_unavailable`, listing the taken seats. Nothing is ever half-held.
- **No deadlocks.** Postgres row locks serialise overlapping holds. Seat IDs are sorted first, so overlapping holds lock rows in the same order.
- **No sweeper needed.** Expired holds are reclaimed by the same predicate.
- **Other rules:**
  - prices are fixed at hold time, with a S$4 booking fee per seat;
  - one active hold per fan per concert;
  - 1–3 seats per order, and at most the ticket type's `max_per_order` (also at most 3, enforced by a CHECK);
  - the sale window and the concert's status.

- **Confirm (simulated payment).** Seats move `held → sold` only if the hold is the fan's and hasn't expired, and the order becomes `confirmed` in the same transaction. Then `OrderConfirmed` is published.
- **Failure paths.** A failed payment (`402`), an expired hold or a cancelled concert releases the seats and closes the order as `cancelled` or `expired`.

`test/seat-inventory.test.ts` races two holds for the same seats against real Postgres and asserts that exactly one wins.

Why Postgres, not DynamoDB, for holds (instructor feedback c): the hold, the order and the seat status must change together. In one database that is a local transaction; across two stores it would need a saga with compensations.

6.4 UC4/UC2: e-ticket pipeline

`ticket-fn` (Node.js 22, arm64, outside the VPC, 256 MB, 60 s) takes up to 10 SQS records per invocation and reports partial batch failures, so only failed records are retried. It handles two events. For each `OrderConfirmed` it:

1. **checks the marker:** stops if `tickets/<orderId>.sent` already exists;
2. **renders the PDF** (pdf-lib): concert, venue, start time in Singapore time, order ID with a QR code, one line per seat, totals. The PDF is written to `tickets/<orderId>.pdf`;
3. **finds the fan's email:** Cognito `AdminGetUser` with the fan's `sub`;
4. **sends** the email with the PDF attached through SES v2;
5. **writes the `.sent` marker**, then logs `deliveryLatencyMs`, which feeds the `ETicketDeliveryLatency` EMF metric (NFR-7).

Idempotency. The marker is written only *after* SES accepts the email, which gives three properties:

- **A redelivery does nothing.**
- **A failure before the send retries the whole record**, so no email is ever lost. Overwriting the PDF is harmless: its content is deterministic.
- **A duplicate email is possible only in two rare cases:** a crash between the send and the marker write, or two copies of one message processed at the same moment. We accept that.

We rejected a conditional `If-None-Match` write of the PDF as the "done" flag, because it would mark the order done before the email was sent.

Failures and download. After 5 receives a record moves to `ticket-dlq` (alarmed). Re-driving it is safe, because orders already sent are skipped. In UC2, `GET /api/me/tickets/:id/download` first calls `HeadObject`:

- **no PDF yet** (just bought, or the event is in the DLQ): `409 ticket_not_ready`, so the fan never lands on an S3 error page;
- **PDF present:** a 300 s pre-signed URL that downloads as an attachment.

Cancellation notices. For each `ConcertCancelled` the same function sends a plain-text email with the reason and the refund amount, guarded by its own marker (`tickets/<orderId>.cancelled.sent`). The event contract and this handling come from Shi Hui's parallel implementation (PR #20), merged into the deployed function.

SES sandbox. SES is in the sandbox (200 emails per 24 h, 1 per second). The sender and every recipient must be verified, which is why the Lambda's maximum concurrency is 2. Production access is a go-live step (Section 12).

6.5 UC5: organiser management, dashboard and reporting

- **Concerts and layout.**
  - Organisers (Cognito group `organisers`, checked by the API) create drafts, replace ticket types, and replace the whole seat layout in one transaction.
  - Publishing needs ticket types and a matching layout.
  - Once the sale opens the layout locks, enforced both in the service (`SELECT … FOR UPDATE`, `409 layout_locked`) and by database triggers.
  - A concert can be cancelled even after the lock. The confirmed orders are read in the cancelling transaction (the concert row lock waits for any in-flight checkout), and after the commit one `ConcertCancelled` per order goes to SNS in batches of 10; `ticket-fn` emails each fan (Section 6.4). If SNS rejects some, cancelling again re-publishes to every order without changing the concert, and `ticket-fn` emails only the fans it missed.
- **Live dashboard.** `GET …/concerts/:id/dashboard` returns available, held and sold counts and revenue per ticket type (`no-store`). The page polls it every 5 s.

- **Reporting views.** `0003_uc5_reporting_views.sql` adds five plain (not materialized) views: `report_confirmed_orders`, `report_sales_by_concert`, `report_sales_by_ticket_type`, `report_sales_by_day` (Singapore calendar day) and `report_seat_status`. Sales count confirmed orders only, money is integer cents per currency, every view carries `organiser_id`, and no view holds buyer identities.
- **Reporting service (instructor feedback e).** A sixth ECS service answers `GET /api/reports/sales?days=7|30|90|all&concertId=`, always filtered by `organiser_id` = the caller, with `Cache-Control: no-store`. It connects as `reporting_ro` (migration `0004`): `SELECT` on the five views only, `default_transaction_read_only = on`, connection limit 10 (2 tasks × `DB_POOL_MAX` 4 = 8). Its password is set inside the VPC through ECS Exec (only a SCRAM verifier crosses) and stored in a Secrets Manager secret that Terraform creates empty, so it is in neither code nor state; ECS injects it at task start. The views are live rather than precomputed, so a sale shows in the reports at once; a read replica is the step for production scale.
- **Sales reports page** (`/organiser/reports`; screenshot `evidence/demo/demo-18-organiser-reports.png`): range and concert filters, KPI cards, revenue by day and by concert, ticket types, seat status, recent orders, CSV download.
- **Tests and live check.** `test/reporting-views.test.ts` checks the views row by row against orders placed through the real services; `test/reporting.test.ts` runs every report query under `SET ROLE reporting_ro` (same rows; reading `concerts` directly is refused, `42501`). On dev, `pg_stat_activity` shows the reporting task as `reporting_ro` and the other services as the master login; an organiser token gets 200, a fan 403, no token 401.

6.6 UC1 and UC2

- **UC1 (public).** Lists published concerts with search, filters and pagination; shows detail with availability per ticket type; serves the seat map. Cache headers: list `public, max-age=60`, detail `max-age=15`, seat map `no-store`.
- **UC2.** Returns the caller's confirmed orders (newest first, at most 100), a single e-ticket, and the download link. Another fan's order answers `404`, so order IDs can't be probed.
- **Auth.** The API verifies every Cognito access token a second time (signature, issuer, expiry, `token_use`, client ID), so a misconfigured gateway route cannot bypass authorisation.

7. Scalability and availability

7.1 Scaling mechanisms

Mechanism	Setting	Effect
CloudFront	SPA cached at the edge (1 day; hashed assets 1 year); each deploy invalidates `/*`	Browsing and the waiting page cost the origin almost nothing
API cache headers	List `max-age=60`, detail `max-age=15`; seat map, checkout and dashboard `no-store`	Caches static data; fast-changing data never cached (Google Doc, system UC4). Browser-only today: the CloudFront `/api/*` behaviour uses `CachingDisabled` (Section 12)
Waiting room	500 admissions per concert per minute; 15-minute admission	Turns a 50,000–100,000-fan spike into a bounded stream of buyers
API Gateway throttle	100 rps steady, 200 burst	Sheds excess with `429` before it reaches the tasks
WAF rate rule	2,000 requests / 5 min per IP	Stops single-source floods and simple bots
ECS autoscaling	Per service: target tracking on CPU 60% and 500 ALB requests per task per minute; dev max 2–3 per service, target design min 2. Measured scale-out (decided / running): 3.9 / 4.2 min and 4.7 / 5.5 min (one service, 1 Nov), 4.7 / 5.6 min (discovery, 6 Nov)	Adds tasks only to the use case under load and removes them afterwards. Too slow alone for a spike that starts from one task (Section 10), so on-sales need pre-scaling (Section 12)
DynamoDB, SQS, SNS, Lambda	On-demand, per request	Nothing to provision on the parts that face the spike
RDS	`db.t4g.micro`; `DB_POOL_MAX` 10 per task (4 for reporting); ≈ 80 `max_connections`	The five master-login services at their maximums (13 tasks) could open 130 connections, above the limit; the connections alarm warns first. Mitigation before production: RDS Proxy in front of the database (pools multiplexed onto few connections), or smaller pools / a larger class; then a read replica for reporting

7.2 Availability: target vs implementation

Failure	Target HA design	Implementation mode	Evidence
F1: task lost	ALB drops it; tasks in the other AZ keep serving; ECS replaces it	ECS replaces the service's only task; that use case is down until it passes the health check	Drill F1 (Section 10)
F2: AZ a lost	Tasks in AZ b serve; RDS fails over (typically 60–120 s, no committed data lost)	Full outage until the AZ recovers or we restore in AZ b	Design only
F3: bad deploy	Circuit breaker rolls back to the last good task definition	Same	ECS events
F4: data corruption	Point-in-time restore to a new instance (7-day window)	Same, 1-day window	Drill F4 (Section 10)
F5: `ticket-fn` failing	Retries, then DLQ + alarm. The order is already confirmed in RDS; a re-drive skips orders already sent	Same	DLQ alarm
F6: region lost	Cross-region snapshot copy and Terraform re-apply (RPO ≤ 24 h, RTO ≤ 4 h)	Not covered	Design only

NFR-5 (99.9%, survive an AZ loss) is therefore a target-design claim. In implementation mode, accepted by the instructor, AZ a is a single point of failure for the API and the database. CloudFront, API Gateway, the ALB, S3, DynamoDB, SQS, SNS and Lambda are regional or global managed services and stay multi-AZ in both modes. The error budget of 99.9% is 43.8 min/month; one task-loss outage (76 s, drill F1) uses ≈ 3% of it.

8. Monitoring and alerts

Monitoring is a Terraform module, `infra/modules/monitoring`: **30 alarms** (plus the messaging module's `ticket-dlq-not-empty`, **31** in total, all OK after the 7 Nov end-to-end pass; `evidence/cloudwatch-alarms.md`), one dashboard, the SNS topic `alerts` with an email subscription to `tayelroy@gmail.com`, and an AWS Budgets budget. Every alarm sends both its ALARM and OK transitions. In the table, *N × 1 min* means N consecutive one-minute datapoints must all breach; *× 6* means one alarm per ECS service.

Signal	Metric	Alarm when	Period	Proposal row / NFR
API latency	ALB `TargetResponseTime` p95	> 300 ms	3 × 1 min	API latency / NFR-1
API errors	ALB target 5xx ÷ `RequestCount`	> 1%	5 min	API errors
Service down (× 6)	`HealthyHostCount` min per target group (missing data = breaching)	< 1	2 × 1 min	Instance health / NFR-5
ECS CPU / memory	`MAX` over the six services of `CPUUtilization`, `MemoryUtilization` avg	> 80%	5 × 1 min	Replaces the EC2 rows
Scaling ceiling (× 6)	ECS `LiveTaskCount` min per service	≥ its max capacity	5 × 1 min	Replaces "ASG at max"
DB CPU	RDS `CPUUtilization`	> 75%	5 × 1 min	Database load
DB connections	RDS `DatabaseConnections` max	> 64 (80% of ≈ 80)	5 × 1 min	Database load
DB storage	RDS `FreeStorageSpace` min	< 2 GiB (10%)	5 min	Database load
Gateway errors	API Gateway `5xx` ÷ `Count`	> 1%	5 min	API errors (also catches VPC link failures)
Gateway latency	API Gateway `Latency` p95	> 500 ms	3 × 1 min	API latency (ALB budget + authorizer + VPC link)
E-ticket backlog	`ticket-queue` `ApproximateAgeOfOldestMessage`	> 60 s	1 min	E-ticket backlog / NFR-7
E-ticket failures	`ticket-dlq` visible messages (messaging module)	> 0	—	E-ticket backlog / NFR-7
E-ticket latency	`TicketMastery/ETicketDeliveryLatency` p95 (EMF)	> 60,000 ms	5 min	NFR-7
Waiting-room DLQ	`join-queue-dlq.fifo` visible messages	> 0	5 min	Fans who never got a position
Lambda errors / throttles	`Errors`, `Throttles` sum, per function (`queue-manager`, `ticket-fn`)	> 0	5 min	Application health
Waiting-room size	`TicketMastery/WaitingRoomSize` max (EMF)	> 20,000	1 min	Waiting room size / NFR-6
Checkout failures	`CheckoutPaymentFailures` ÷ `CheckoutPaymentAttempts` (log metric filters on the confirm route: 200 vs 402)	> 20% with ≥ 5 attempts	5 min	Checkout failures
Budget	AWS Budgets, monthly, USD 80	Actual > USD 20, 40, 60, 80; forecast > 80	daily	Budget / NFR-11

Proposal rows adapted. Every proposal row has an alarm except three. EC2 status checks and the ASG row became the ECS and target-health alarms. The ElastiCache row was dropped, because holds are Postgres rows covered by the RDS alarms. The waiting-room metric is per minute rather than every 10 s, because one minute is the Scheduler's shortest interval.

Two metrics are produced inside our own code:

- **Waiting-room size.** `queue-manager` has no NAT, so it cannot call `PutMetricData`. It prints one EMF JSON line per run, and CloudWatch Logs turns it into a metric.
- **Checkout failures.** Counted from the API's request log by two metric filters, with no extra code.

Dashboard `ticketmastery-dev`: an alarm-status panel plus 19 graphs, each with its threshold line: ALB latency and 5xx rate; requests, healthy and unhealthy targets per service; ECS CPU, memory and live tasks per service (one line each); RDS CPU and credits, connections, free storage; API Gateway requests, 5xx rate and latency; `ticket-queue` age and depth, e-ticket delivery latency; both DLQs; Lambda errors and throttles; waiting-room size; checkout failures. Per-service CPU, memory and unhealthy-host alarms were dropped as noise (a first split version had 41 alarms).

Alerts in practice. During NFR-2 run 1 (Section 10) the alarms did their job without being forced: the API Gateway 5xx-rate alarm fired 2 min 40 s after the ramp started (about a minute after the first failed requests), followed by API Gateway p95 latency, ALB 5xx rate and ALB p95 latency; all four returned to OK ≈ 3–8 minutes after the scale-out decision (ALB 5xx last, at 03:17:49), and `ecs-at-max-capacity` then fired because the service sat at its maximum. The F1 drill raised the API Gateway 5xx alarm again, and that email arrived. For an end-to-end check we also forced `ticketmastery-dev-join-dlq-not-empty` into `ALARM` with `aws cloudwatch set-alarm-state`: SNS delivered the ALARM (09:29:54) and OK (09:30:40) emails 46 s apart, with 0 failed notifications. (The email subscription was confirmed only after the burst runs, so those alarms reached SNS but not the inbox; `evidence/monitoring-alerts.md`.) In NFR-2 run 3, on the split stack, only `discovery-ecs-at-max-capacity` fired, naming the one service under load.

Runbook. Alerts go to one team inbox today; a 429/throttle alarm and an on-call rota are future work.

Alarm	First action	Owner
Service down, ALB/API Gateway 5xx or p95	ECS events and target health; roll back the task definition if a deploy is in flight	Elroy
`ecs-at-max-capacity`, DB connections/ CPU	Raise the service's min/max in `tfvars` (targeted apply); check `pg_stat_activity`	Elroy
`ticket-dlq-not-empty`, e-ticket age/latency	Read `ticket-fn` logs, fix, re-drive the DLQ (already-sent orders are skipped)	Shi Hui
`join-queue` DLQ, waiting-room size	Read `queue-manager` logs; re-drive; check the admit schedule is enabled	Yan Ting

9. Management approach

9.1 Infrastructure as code and remote state

All infrastructure is Terraform (≥ 1.10, AWS provider ~> 5.0). One root module (`infra/envs/dev`) wires one module per area, each owned by one person:

Module	Owner
`network`, `security`, `ecr`, `rds`, `ecs-service`, `github-deploy`, `monitoring`	Elroy
`edge` (and the `deploy-frontend.yml` workflow)	Cody
`auth-apigw`	Andrew
`waiting-room`	Yan Ting
`messaging`	Shi Hui

Conventions every module follows:

- shared values come in as variables, never data-source lookups;
- names are prefixed `ticketmastery-dev`;
- the module that owns a resource grants the API least-privilege access to it;
- non-secret API configuration goes in `extra_environment`;
- secrets are read from Secrets Manager at runtime.

State and safety. State lives in S3 (`ticketmastery-tfstate-000246635159`, versioned) with S3-native locking, so only one apply runs at a time. One `tfvars` block switches between the implementation and HA modes. Preconditions on the RDS, ECS and API Gateway resources stop a run on invalid settings; `check` blocks warn (without failing) on risky combinations such as private-subnet tasks without egress or a VPC link to a public ALB. Every change is `plan -out` → review → `apply` of that plan, announced to the team first.

9.2 CI/CD and operations

- **CI** (`ci.yml`, on pull requests and pushes to `main`): API type-check and tests against a Postgres 16 service with the coverage gate (100% for `src/auth`, `src/config.ts`, `src/modules/account`), `ticket-fn` tests plus a committed-bundle check, `queue-manager` unittests, and `terraform fmt -check` + `validate`; all green on 7 Nov. `frontend.yml` type-checks and builds the frontend.
- **Frontend deploy** (`deploy-frontend.yml`) is manual and runs only from `main`. It builds the SPA, syncs it to S3 and invalidates CloudFront. It authenticates through GitHub OIDC, so no AWS keys are stored in GitHub.

- **API image:** built for `linux/amd64`, tagged with the git SHA (`4a443fa` on all six services), pushed to ECR (scanned on push), and rolled out by hand with Terraform; the deployment circuit breaker rolls back failures.
- **Incident: cutover 503.** Moving from one service to six gave ≈ 5 min of 503 on dev. Cause: the old target group was deleted before the listener switched. Fix: a targeted apply. Prevention: create new target groups and rules before removing the old ones; later deploys are rolling.
- **Tests** run against real Postgres (seat-hold race, layout triggers, reporting views and role, auth, admission).
- **ECS Exec** (IAM-gated) runs one-off checks inside the VPC, such as the oversell SQL, without a bastion.

9.3 Branching, ownership and review

- **Branches.** Trunk-based on `main` with short-lived feature branches and Conventional Commit titles.
- **Ownership.** `.github/CODEOWNERS` maps each module, migration use-case tag and Lambda to its owner. Every PR needs approval from the owner of each path it touches.
- **PR template.** Covers type, touched areas and how the change was tested, plus a checklist: rebased; no secrets, `tfvars` or state committed; no edits to merged migrations; `terraform fmt/validate` pass; team told before any apply; docs updated.
- **Merging.** Feature PRs are squash-merged. The final integration branch is merged with a merge commit so each teammate keeps authorship.

9.4 Cost management

- **Implementation mode by default.** No NAT, one AZ and an internal ALB took the base stack from ≈ USD 110–115 (one NAT, public ALB) to ≈ 62/month; the six-service split (one task each) raised it to ≈ 136 (whole stack ≈ 150).
- **Teardown.** `terraform destroy` after demos and work sessions; recreating the stack takes about 15 minutes.
- **Budgets.** Alerts on actual spend at USD 20/40/60/80, and on a forecast above USD 80 (NFR-11). At ≈ 4.9/day a full month would pass USD 80, so the forecast alert will fire; actual November spend stays under it because the stack is destroyed after 10 Nov: ≈ 6 days × 2.5 + ≈ 4 days × 4.9 ≈ **USD 35** (estimate). Total project spend (NFR-11): October 22.92 + November ≈ 35 ≈ **USD 58** (estimate), under 80.
- **Schedule.** The admit schedule costs ≈ USD 0 always on (free tiers). It can still be paused, or limited to start and end dates (instructor feedback d).
- **Pay per request** for everything idle between on-sales: waiting room, e-tickets and API Gateway.

Actual spend. Cost Explorer shows **USD 22.92** of usage in October 2026 (stack deployed from 22 Oct), all covered by credits, so USD 0 was charged. A full day cost USD 2.02 after the ALB became internal (estimate then ≈ 2.5). The largest items match the estimate's order: ALB 8.26, RDS 6.21, Fargate 3.34, public IPv4 2.93, WAF 2.05 (`evidence/nfr11-cost.md`, Appendix B).

9.5 Security

Layer	Control
Edge	HTTPS only (CloudFront redirect); WAF common rules + IP rate limit; S3 reachable only through OAC
Identity	Cognito with email verification; organisers as a group; JWT checked by API Gateway and again in the API
Network	Internal ALB reachable only from the VPC link SG; SG chain → task → RDS; RDS private, not publicly accessible
Data	RDS storage encrypted; TLS to the DB; e-ticket bucket private (public access blocked, SSE-S3, HTTPS-only bucket policy), read through 300 s pre-signed URLs; the reporting service logs in as `reporting_ro` (read-only, the five views only)
Secrets	Master password generated by RDS in Secrets Manager; `reporting_ro` password set inside the VPC through ECS Exec into a secret Terraform creates empty; both injected as ECS secrets, never in code, `tfvars` or state
Ops	GitHub OIDC role limited to `main`, the frontend bucket and the distribution; ECS Exec gated by IAM
App	zod validation of all input; errors never expose internals; dev-header auth refused in production

Gaps accepted for dev (production steps in Section 12):

- **API Gateway → ALB → task is plain HTTP on port 80.** The hop never leaves the VPC and only the VPC link can reach the ALB, but it is not encrypted. HTTPS on the ALB needs a certificate and a custom domain.
- **Cognito MFA is off.**
- **One task role is shared by all six services,** so each service can reach the queues and buckets of the others.
- **Deletion protection is off,** so `terraform destroy` works.
- **SES is in the sandbox.**

- **DB TLS without server-certificate verification** (`pool.ts`): encrypted, but not protected against an impersonated endpoint inside the VPC; production loads the RDS CA bundle.

9.6 Backup, disaster recovery, RPO and RTO

Scenario	Mechanism	RPO (NFR-8)	RTO (NFR-9)	Status
Task loss	ECS replaces the service's task; ALB health check	n/a	≤ 5 min	Drill F1
AZ loss	Multi-AZ RDS synchronous standby; tasks in AZ b	0	≤ 5 min	Design only
Data corruption	Automated backups + point-in-time restore to a new instance	≤ 30 min	≤ 1 h	Drill F4: measured RPO 2.8 min, RTO 9 min 14 s (DB accepting connections)
Region loss	Cross-region snapshot copy + Terraform apply in another region	≤ 24 h	≤ 4 h	Design only
Lost e-ticket	The order in RDS is the source of truth; a DLQ re-drive regenerates the e-ticket	0	minutes	DLQ alarm

Infrastructure needs no backup beyond the repository and the versioned state bucket, because every resource is rebuilt from code.

9.7 Teardown

`terraform destroy` removes the stack (`docs/deploy.md`). Dev defaults keep it clean: no final snapshot, no deletion protection, ECR and the e-ticket bucket force-deleted. Afterwards we check for anything that keeps billing:

- leftover ENIs from the in-VPC Lambda;
- unattached Elastic IPs;
- log groups created outside Terraform;
- manual snapshots.

The state bucket and the budget are kept.

10. Evidence

Method. Repeatable scripts in `loadtest/` (README has every command): k6 drives the API; shell scripts read CloudWatch, ECS and RDS; two drills change things (F1 stops the task, F4 restores RDS to a new instance and deletes it). Each run writes `loadtest/results/<test>-<UTC time>.{json,md}`; the curated copies are in `evidence/`. Conditions: tests call the **API Gateway URL directly**, bypassing CloudFront and the WAF (its 2,000 requests per 5 min per IP would block one test machine); the stage throttle (100 req/s, burst 200) still applies, and `429`s are counted apart from failures; 125 Cognito test users with real tokens and SES mailbox-simulator addresses; k6 runs from Singapore, and the ALB's `TargetResponseTime` gives the server-side share.

NFR	Method	Result (on the deployed stack, 31 Oct–1 Nov 2026 unless marked 6 Nov)	Pass?
NFR-1	10 req/s for 5 min of list / detail / seat map	3,000 requests, 0 errors; k6 **p95 30 ms, p99 43 ms** (seat map p95 35 ms); ALB p95 11 ms, p99 20 ms; CPU ≤ 15.4%	Pass
NFR-2	10 ⟶ 100 req/s (10×) in 1 min, held 10 min. Run 1: 1 task, max 3 (implementation). Run 2: min 2, max 4 (the target design's two-task floor)	Run 1: **2,153 of 64,964 requests failed (3.3%)** in the 4 min before scale-out: one 0.25 vCPU task hit 94% CPU at ≈ 60 req/s, its `/health` failed (503) and ECS replaced it; 1 ⟶ 3 tasks decided at 3.9 min, running at 4.2 min, then 0 errors. Run 2: **0 of 69,601 failed**, 0 throttled, k6 p95 31 ms, CPU ≤ 54%; 2 ⟶ 4 tasks decided at 4.7 min, running at 5.5 min (Figure 4)	Run 1 fail; run 2 pass on errors, 5.5 vs 5 min scale-out
NFR-2 run 3 (6 Nov)	Same 10 ⟶ 100 req/s ramp with an 8-min peak, on the six-service stack, 1 task per service	57,600 requests, **0 dropped, 0 throttled**; k6 p95 63 ms (peak phase 71 ms), p99 179 ms. **Only discovery scaled**: 1 ⟶ 3 tasks decided at 4.7 min, running at 5.6 min; the other four API services stayed at 1 task	Pass on errors; 5.6 vs 5 min scale-out
NFR-3	100 admitted fans race for 20 seats (2 per hold, 10% of payments declined), twice; then SQL inside the VPC (ECS Exec)	299 hold attempts, **278 lost races (409)**, 20 orders for 40 seats, 1 declined seat re-sold, 0 server errors. SQL: **0 seats in two confirmed orders**, 0 sold seats without an order, per concert and for the whole database	Pass
NFR-4	Hold a seat, poll the seat map every 2 s around expiry, re-hold, pay late	Available **0.6 s** after expiry; fan B re-held it (`201`); fan A's late payment `409 hold_expired`	Pass (≤ 30 s)
NFR-5	F1: `ecs stop-task` on the only task (single-service stack); `/health` polled every second	Replacement started at +54 s, healthy at +82 s; **outage 76 s**. With two tasks the ALB routes around a lost task; AZ loss is design only	Pass (≤ 5 min)
NFR-6	20 fans join ≈ 236 ms apart on average before the sale; admit batch lowered to 5 for this run, then restored to 500	Positions 1–20 = join order; **4 admit runs** one minute apart; **0 inversions** in position or admission	Pass
NFR-7	`ticket-fn`'s `deliveryLatencyMs` (SES accepted – order confirmed) over the NFR-3 orders	Run 2: **10/10 emailed within 60 s, p95 9.9 s**. Run 1 exposed an IAM gap (SESv2 raw sends need `ses:SendRawEmail`): all 10 failed once, stayed queued, and were emailed on the retry after the fix; none lost	Pass (after the fix)
NFR-8/9	F4: point-in-time restore to a new private instance, timed from the RDS event log, then deleted	Restore point **168 s** before the request; database accepting connections after **9 min 14 s** (`available` at 28 min)	Pass
NFR-10	Read-only configuration checks	RDS private and KMS-encrypted; DB subnets without an internet route; ALB internal; CloudFront and API Gateway HTTPS only; WAF on; S3 and SQS encrypted. VPC link ⟶ ALB is HTTP inside the VPC (Section 9.5)	Pass
NFR-11	Cost Explorer and Budgets	October usage USD 22.92 (USD 0 after credits); budget USD 80 with alerts. November ≈ USD 35, total ≈ 58 (estimates, stack destroyed after 10 Nov); a full month at ≈ 4.9/day would pass 80	Pass so far; forecast alert expected

Evidence files: `evidence/nfr1-browse.md`, `nfr2-burst-run1.md`, `nfr2-burst-run2.md`, `nfr3-oversell.md`, `nfr4-hold-expiry.md`, `nfr5-f1-task-loss.md`, `nfr6-waiting-room.md`, `nfr7-eticket.md`, `nfr8-f4-pitr.md`, `nfr10-config.md`, `nfr11-cost.md`, `monitoring-alerts.md`, `cloudwatch-alarms.md`, `nfr2-burst-run3-split.md`, `nfr2-burst-run3-split-ecs.md` and `nfr2-burst-run3-tasks-per-service.png` (tasks per service during run 3).

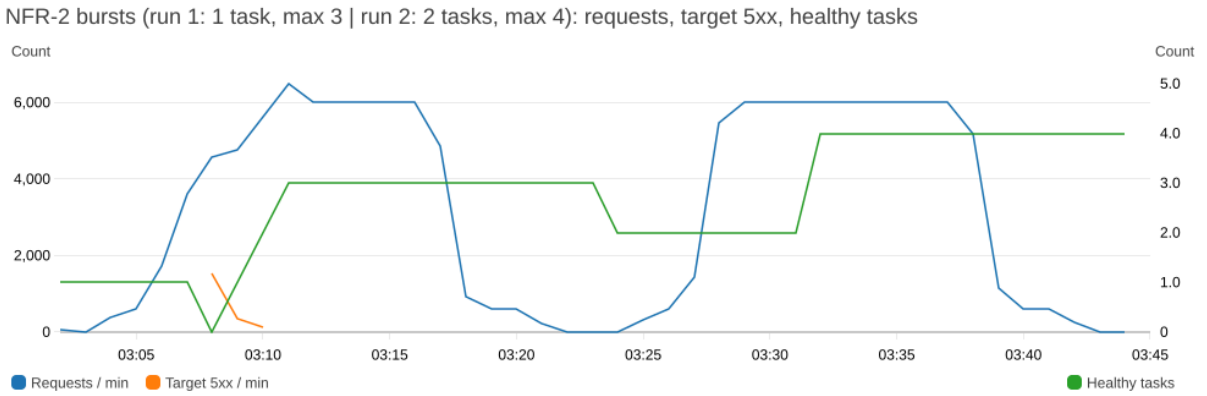


Figure 4: NFR-2 bursts: requests and target 5xx per minute (left axis) and healthy tasks (right axis). Run 1 (03:06–03:17 UTC) starts with one task; run 2 (03:27–03:38 UTC) with two.

What NFR-2 taught us. One 0.25 vCPU task saturates at about 60 req/s. Target tracking needs three one-minute datapoints above target before it acts, and a new task needs about 40 s to pass its health check, so scale-out takes 4–6 minutes. Run 1 also showed a failure mode: the `/health` check includes a database `SELECT 1`, so when the only task was overloaded its health check failed (503) and ECS replaced it, turning slow responses into errors. With two tasks as the floor (the target design), the same burst caused no errors. Run 3, on the split stack, scaled only `discovery`, the service that serves browsing, which is the point of per-feature services. Scale-out still took longer than 5 min, so the fixes for production stay, in order: scale out ahead of each announced on-sale with a scheduled action (the proposal already planned this), keep at least two tasks per service, and add a faster step-scaling policy on a high-resolution alarm. The waiting room bounds the transactional load, but browse traffic is not gated by it.

Functional evidence. On 1 Nov (09:39–09:44 UTC) one fan account went through the whole flow on the CloudFront URL, and an organiser opened the live dashboard (Figure 5; full screenshots in `evidence/screens/`):

1. UC1: concert list and detail (`uc1-concert-list.jpg`, `uc1-concert-detail.jpg`);
2. UC3: queue number #4, admitted 18 s later by the next scheduled run (`uc3-waiting-room-queued.jpg`, `uc3-waiting-room-admitted.jpg`);
3. UC4: two seats selected, the 10-minute hold, simulated payment, confirmation (`uc4-seat-selection.jpg`, `uc4-hold-countdown.jpg`, `uc4-confirmation.jpg`);
4. UC2: My Tickets, the ticket page, and the PDF downloaded through its pre-signed link: 3,878 bytes from the private bucket's `tickets/<orderId>.pdf` (`uc2-my-tickets.jpg`, `uc2-ticket-page.jpg`, `uc2-eticket-pdf.png`). The fan's address is an SES mailbox-simulator address, so the email itself was accepted by SES rather than read in an inbox; the live demo uses a verified inbox;
5. UC5: the organiser's concerts and live dashboard (sold 20/1000, revenue SGD 1,000 after the NFR-3 run) (`uc5-organiser-concerts.jpg`, `uc5-organiser-dashboard.jpg`);
6. UC5 reporting: the views on the live database return the same totals (`evidence/uc5-reporting-views.md`); the Sales reports page on the reporting service, 6 Nov (`evidence/demo/demo-18-organiser-reports.png`);
7. 7 Nov, six-service stack, real browser on CloudFront, on a 10-seat test concert (`evidence/e2e-7nov.md`): sign-in → waiting room (admitted) → 2-seat hold → payment → e-ticket emailed 3.4 s after payment (`deliveryLatencyMs` 3429), PDF downloaded → dashboard "Sold 2 / 10" → Sales reports (SGD 108, 2 tickets, 1 order) → cancel → notice sent (1 order, 1 published); all 31 alarms OK afterwards.

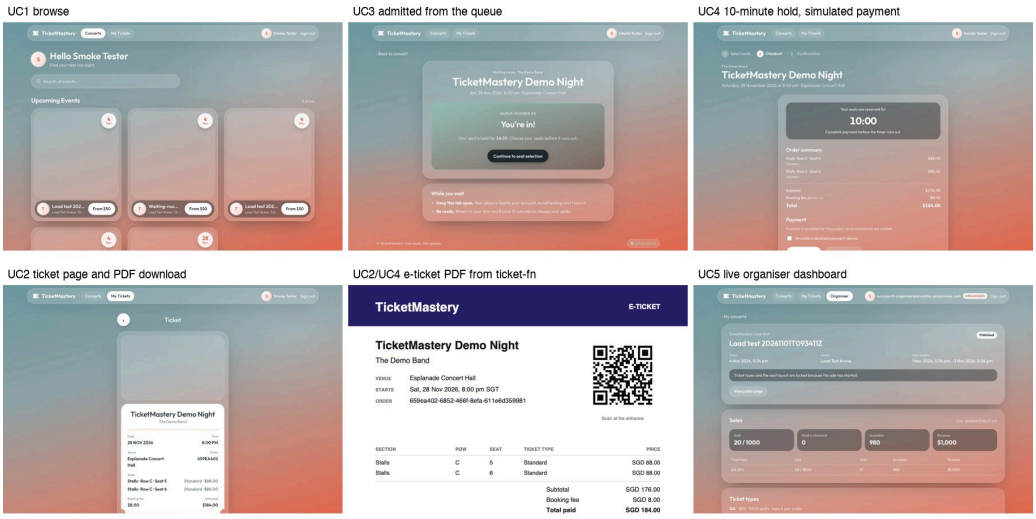


Figure 5: The deployed system on the CloudFront URL (1 Nov 2026).

11. Instructor feedback and our response

#	Feedback at the proposal review	Our response	Status
—	Use public subnets to avoid NAT, one AZ, no HA in the implementation, but show HA in the diagram	Implementation mode is the Terraform default; the HA design sits behind variables and is shown with its cost (Figure 1, Section 3.3)	Done
a	Put API Gateway in front of Fargate to manage all APIs centrally	HTTP API with a Cognito JWT authorizer, throttling, access logs, and a VPC link to an internal ALB (chosen over a public ALB with a secret header)	Done
b	Consider per-feature containers so each part scales separately	Six ECS services from one image (one per use case plus reporting), each with its own target group, ALB path rule, autoscaling policy and life cycle; per-UC migrations, no cross-module imports. NFR-2 run 3 scaled only `discovery` (Section 10)	Done
c	DynamoDB is the usual store for dynamic session data	The live queue and admissions are in DynamoDB. Seat holds stay in Postgres on purpose: one atomic `UPDATE` beside the seats and orders guarantees no double-booking (Section 6.3)	Done, justified
d	EventBridge every minute costs money; check it	≈ 43,800 runs/month: inside the Scheduler (14 M) and Lambda (1 M) free tiers, ≈ USD 0. The schedule can be paused or date-limited, and admits only inside each concert's sale window	Done
e	UC5 needs internal reporting, e.g. QuickSight	Live dashboard; five tested SQL views; a read-only reporting service (`reporting_ro`) and a Sales reports page (Section 6.5). QuickSight was built in Terraform, then removed: the AWS Free-plan account can't subscribe to it	Done (no BI tool)

12. Limitations and future work

- **No BI tool.** QuickSight was built in Terraform and then removed: the account is on the AWS Free plan, and subscribing fails until it moves to a paid plan. The reporting service and Sales reports page cover the need (Section 6.5); an RDS read replica for reporting is the step at production scale (`docs/reporting.md`).
- **HA not deployed,** because the instructor asked for the HA design to be shown, not built. AZ-loss tolerance, Multi-AZ RDS and two tasks per service are design-only; turning them on is a `tfvars` change (≈ USD 310/month). Region-loss and AZ-loss recovery were never exercised.
- **Cost of the split.** Six services at one task each raise the whole stack from ≈ USD 76 (62 base + 13.7 rest) to ≈ 150/month (Appendix B), so a full month would pass the USD 80 budget; the stack is destroyed after 10 Nov.
- **Shared task role.** All six services use one task role (Section 5.3). Next step: one least-privilege task role per service.
- **Faster scaling.** Scale-out took 5.5–5.6 min, over the 5 min target (NFR-2 runs 2 and 3). Pre-scale with an `aws_appautoscaling_scheduled_action` per service that raises the minimum (e.g. to 3) 15 min before `sale_starts_at`; add step scaling on a high-resolution alarm, and use a liveness check that doesn't fail a busy task (NFR-2 run 1, Section 10).
- **Next steps:** CloudFront caching of `GET /api/concerts*`; an adaptive admit batch; Stripe payments (idempotency key, signed webhook); cross-region snapshot copy for region-loss DR (multi-region writes are not justified with 80%

of customers in Singapore); SES production access with DKIM, HTTPS on the ALB, Cognito MFA, deletion protection and 7-day backups, WAF Bot Control.
- **API rollout is manual.** CI tests the API, Lambda and Terraform, but the image is built, pushed and rolled out to the six services by hand (Terraform apply); per-service task roles are also still to do.

Appendix A. API list

Method and path	Auth	Purpose
`GET /health`	public	Readiness (DB `SELECT 1`), used by the ALB
`GET /api/concerts`, `/:id`, `/:id/seats`	public	List and search; detail with availability; seat map
`GET /api/me`, `/api/me/orders`, `/api/me/tickets/:id`, `…/download`	fan	Identity; orders; one e-ticket; pre-signed PDF URL (300 s), `409 ticket_not_ready` until it exists
`POST /api/waiting-room/:concertId/join`, `GET …/status`	fan	Join the queue; position and state
`GET /api/checkout/concerts/:concertId/seats`; `POST`, `DELETE /api/checkout/holds[/:holdId]`	fan (hold: admitted)	Seat map with own hold; hold 1–3 seats for 10 min (`403 not_admitted`, `409 seats_unavailable`); release
`POST /api/checkout/holds/:holdId/confirm`	fan	Simulated payment; `402 payment_failed`, `409 hold_expired`
`POST`, `GET /api/organiser/concerts[/:id]`; `PUT …/ticket-types`, `…/layout`; `POST …/publish`, `…/cancel`; `GET …/dashboard`	organiser	Drafts, lists, detail; replace types and layout; publish, cancel with reason; live counts and revenue
`GET /api/reports/sales` (`days` = 7, 30, 90 or `all`; optional `concertId`)	organiser	Sales reports over the `report_*` views (`reporting_ro`)

Errors are JSON `{ "error": { "code", "message", "details?" } }`: 400 validation, 401 no user, 403 not allowed or not admitted, 404 unknown, 409 conflict.

Appendix B. Cost tables

Source `docs/cost-estimate.md`: on-demand list prices, USD, ap-southeast-1, 730 h/month, excluding tax; estimates. Single-service stack: ≈ 76/month estimated (62 base + 13.7 rest), USD 2.02–2.26/day measured; October USD 22.92, all credits (`evidence/nfr11-cost.md`).

Item	Implementation (deployed)	≈ USD	Target HA (2 tasks per service)	≈ USD
NAT gateways + Elastic IPs	none	0	2 + 2 EIPs	86.5 + 7.3
ALB (internal)	1, ~1 LCU	24.2	1	24.2
RDS `db.t4g.micro`	single-AZ	18.3	Multi-AZ	36.5
RDS storage 20 GB gp3	single-AZ	2.8	Multi-AZ	5.5
Fargate 0.25 vCPU / 0.5 GB (11.2 per task)	6 services × 1 task	67.2	6 × 2 = 12 tasks, private	134.4
Public IPv4 (3.65 each)	6 tasks	21.9	none	0
Secrets Manager (2 secrets) 0.8, Logs 0.9, ECR 0.1		1.8		1.8
Base stack		**136.2** (4.5/day)		**296.2** (9.7/day)
WAF (1 ACL, 2 rules)		7		7
Monitoring: 34 billable alarm metrics (3.4) + dashboard over 50 metrics (3); pay-per-request ≈ 0.3		6.7		6.7
Whole stack		**≈ 150** (149.9; 4.9/day)		**≈ 310** (309.9; 10.2/day)